

weather-prediction-fi

December 9, 2024

1 Requirements and Imports

```
[1]: pip install -r requirements.txt
```

```
Requirement already satisfied: numpy in /opt/conda/lib/python3.11/site-packages  
(from -r requirements.txt (line 1)) (1.24.4)  
Requirement already satisfied: pandas in /opt/conda/lib/python3.11/site-packages  
(from -r requirements.txt (line 2)) (2.0.3)  
Requirement already satisfied: matplotlib in /opt/conda/lib/python3.11/site-  
packages (from -r requirements.txt (line 3)) (3.8.0)  
Requirement already satisfied: seaborn in /opt/conda/lib/python3.11/site-  
packages (from -r requirements.txt (line 4)) (0.13.0)  
Requirement already satisfied: scikit-learn in /opt/conda/lib/python3.11/site-  
packages (from -r requirements.txt (line 5)) (1.3.1)  
Requirement already satisfied: scipy in /opt/conda/lib/python3.11/site-packages  
(from -r requirements.txt (line 6)) (1.11.3)  
Requirement already satisfied: umap-learn in /opt/conda/lib/python3.11/site-  
packages (from -r requirements.txt (line 7)) (0.5.7)  
Requirement already satisfied: umap in /opt/conda/lib/python3.11/site-packages  
(from -r requirements.txt (line 8)) (0.1.1)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
/opt/conda/lib/python3.11/site-packages (from pandas->-r requirements.txt (line  
2)) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.11/site-  
packages (from pandas->-r requirements.txt (line 2)) (2023.3.post1)  
Requirement already satisfied: tzdata>=2022.1 in /opt/conda/lib/python3.11/site-  
packages (from pandas->-r requirements.txt (line 2)) (2023.3)  
Requirement already satisfied: contourpy>=1.0.1 in  
/opt/conda/lib/python3.11/site-packages (from matplotlib->-r requirements.txt  
(line 3)) (1.1.1)  
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.11/site-  
packages (from matplotlib->-r requirements.txt (line 3)) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in  
/opt/conda/lib/python3.11/site-packages (from matplotlib->-r requirements.txt  
(line 3)) (4.43.1)  
Requirement already satisfied: kiwisolver>=1.0.1 in  
/opt/conda/lib/python3.11/site-packages (from matplotlib->-r requirements.txt  
(line 3)) (1.4.5)
```

```
Requirement already satisfied: packaging>=20.0 in
/opt/conda/lib/python3.11/site-packages (from matplotlib->-r requirements.txt
(line 3)) (23.2)
Requirement already satisfied: pillow>=6.2.0 in /opt/conda/lib/python3.11/site-
packages (from matplotlib->-r requirements.txt (line 3)) (10.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/opt/conda/lib/python3.11/site-packages (from matplotlib->-r requirements.txt
(line 3)) (3.1.1)
Requirement already satisfied: joblib>=1.1.1 in /opt/conda/lib/python3.11/site-
packages (from scikit-learn->-r requirements.txt (line 5)) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/opt/conda/lib/python3.11/site-packages (from scikit-learn->-r requirements.txt
(line 5)) (3.2.0)
Requirement already satisfied: numba>=0.51.2 in /opt/conda/lib/python3.11/site-
packages (from umap-learn->-r requirements.txt (line 7)) (0.57.1)
Requirement already satisfied: pynndescent>=0.5 in
/opt/conda/lib/python3.11/site-packages (from umap-learn->-r requirements.txt
(line 7)) (0.5.13)
Requirement already satisfied: tqdm in /opt/conda/lib/python3.11/site-packages
(from umap-learn->-r requirements.txt (line 7)) (4.66.1)
Requirement already satisfied: llvmlite<0.41,>=0.40.0dev0 in
/opt/conda/lib/python3.11/site-packages (from numba>=0.51.2->umap-learn->-r
requirements.txt (line 7)) (0.40.1)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.11/site-
packages (from python-dateutil>=2.8.2->pandas->-r requirements.txt (line 2))
(1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
[2]: import pyspark
print(pyspark.__version__)
```

3.5.0

```
[3]: import numpy as np
import pandas as pd
from scipy import stats
from scipy.sparse import hstack
import matplotlib.pyplot as plt
import seaborn as sns

from pyspark.sql import SparkSession
from pyspark.ml.feature import VectorAssembler, StringIndexer

from pyspark.ml.classification import RandomForestClassifier

from pyspark.ml.evaluation import BinaryClassificationEvaluator
```

```

from pyspark.ml.classification import GBTClassifier
from pyspark.ml import Pipeline

from pyspark.ml.classification import LogisticRegression


from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

import umap
from umap import UMAP

import warnings

```

2 Data Summary

```
[4]: file_path = 'weatherAUS.csv'
try:
    df = pd.read_csv(file_path)
    print(f"File '{file_path}' loaded successfully.")
except FileNotFoundError:
    print(f"Error: The file '{file_path}' was not found. Please check the file path and try again.")
df = None

total_size = df.size
print("Total number of elements:", total_size)
```

File 'weatherAUS.csv' loaded successfully.

Total number of elements: 3345580

```
[5]: print("Preview of the Dataset (Head):")
display(df.head())
```

Preview of the Dataset (Head):

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	

```

4 2008-12-05 Albury 17.5 32.3 1.0 NaN NaN
   WindGustDir WindGustSpeed WindDir9am ... Humidity9am Humidity3pm \
0 W 44.0 W ... 71.0 22.0
1 WNW 44.0 NNW ... 44.0 25.0
2 WSW 46.0 W ... 38.0 30.0
3 NE 24.0 SE ... 45.0 16.0
4 W 41.0 ENE ... 82.0 33.0

Pressure9am Pressure3pm Cloud9am Cloud3pm Temp9am Temp3pm RainToday \
0 1007.7 1007.1 8.0 NaN 16.9 21.8 No
1 1010.6 1007.8 NaN NaN 17.2 24.3 No
2 1007.6 1008.7 NaN 2.0 21.0 23.2 No
3 1017.6 1012.8 NaN NaN 18.1 26.5 No
4 1010.8 1006.0 7.0 8.0 17.8 29.7 No

RainTomorrow
0 No
1 No
2 No
3 No
4 No

```

[5 rows x 23 columns]

```
[6]: print("\nSummary Statistics for Numerical Columns:")
numerical_summary = df.describe(include=[float, int])
display(numerical_summary)
```

Summary Statistics for Numerical Columns:

	MinTemp	MaxTemp	Rainfall	Evaporation	\
count	143975.000000	144199.000000	142199.000000	82670.000000	
mean	12.194034	23.221348	2.360918	5.468232	
std	6.398495	7.119049	8.478060	4.193704	
min	-8.500000	-4.800000	0.000000	0.000000	
25%	7.600000	17.900000	0.000000	2.600000	
50%	12.000000	22.600000	0.000000	4.800000	
75%	16.900000	28.200000	0.800000	7.400000	
max	33.900000	48.100000	371.000000	145.000000	
	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	\
count	75625.000000	135197.000000	143693.000000	142398.000000	
mean	7.611178	40.035230	14.043426	18.662657	
std	3.785483	13.607062	8.915375	8.809800	
min	0.000000	6.000000	0.000000	0.000000	
25%	4.800000	31.000000	7.000000	13.000000	

50%	8.400000	39.000000	13.000000	19.000000	
75%	10.600000	48.000000	19.000000	24.000000	
max	14.500000	135.000000	130.000000	87.000000	
	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm	\
count	142806.000000	140953.000000	130395.000000	130432.000000	
mean	68.880831	51.539116	1017.64994	1015.255889	
std	19.029164	20.795902	7.10653	7.037414	
min	0.000000	0.000000	980.50000	977.100000	
25%	57.000000	37.000000	1012.90000	1010.400000	
50%	70.000000	52.000000	1017.60000	1015.200000	
75%	83.000000	66.000000	1022.40000	1020.000000	
max	100.000000	100.000000	1041.00000	1039.600000	
	Cloud9am	Cloud3pm	Temp9am	Temp3pm	
count	89572.000000	86102.000000	143693.000000	141851.000000	
mean	4.447461	4.509930	16.990631	21.68339	
std	2.887159	2.720357	6.488753	6.93665	
min	0.000000	0.000000	-7.200000	-5.40000	
25%	1.000000	2.000000	12.300000	16.60000	
50%	5.000000	5.000000	16.700000	21.10000	
75%	7.000000	7.000000	21.600000	26.40000	
max	9.000000	9.000000	40.200000	46.70000	

```
[7]: print("\nSummary Statistics for Categorical Columns:")
categorical_summary = df.describe(include=[object])
display(categorical_summary)
```

Summary Statistics for Categorical Columns:

	Date	Location	WindGustDir	WindDir9am	WindDir3pm	RainToday	\
count	145460	145460	135134	134894	141232	142199	
unique	3436	49	16	16	16	2	
top	2013-11-12	Canberra	W	N	SE	No	
freq	49	3436	9915	11758	10838	110319	
	RainTomorrow						
count		142193					
unique		2					
top		No					
freq		110316					

```
[8]: print("\nNull Values in Each Column:")
null_values = df.isnull().sum()
display(null_values)
```

Null Values in Each Column:

```
Date          0
Location      0
MinTemp       1485
MaxTemp       1261
Rainfall       3261
Evaporation   62790
Sunshine       69835
WindGustDir   10326
WindGustSpeed 10263
WindDir9am    10566
WindDir3pm    4228
WindSpeed9am  1767
WindSpeed3pm  3062
Humidity9am   2654
Humidity3pm   4507
Pressure9am   15065
Pressure3pm   15028
Cloud9am      55888
Cloud3pm      59358
Temp9am       1767
Temp3pm       3609
RainToday      3261
RainTomorrow   3267
dtype: int64
```

3 Missing Values

```
[9]: binary_columns = ['RainToday', 'RainTomorrow'] # only for yes or no values
for col in binary_columns:
    if col in df.columns:
        mode_value = df[col].mode()[0]
        df[col].fillna(mode_value, inplace=True)
        print(f"Filled missing values in '{col}' with the mode: {mode_value}")
```

Filled missing values in 'RainToday' with the mode: No
Filled missing values in 'RainTomorrow' with the mode: No

```
[10]: # Fill missing values in numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
for col in numerical_columns:
    if df[col].isnull().sum() > 0:
        mean_value = df[col].mean() # can use median
        df[col].fillna(mean_value, inplace=True)
        print(f"Filled missing values in numerical column '{col}' with the mean:
              ↵ {mean_value}")
```

```

Filled missing values in numerical column 'MinTemp' with the mean:
12.19403438096892
Filled missing values in numerical column 'MaxTemp' with the mean:
23.22134827564685
Filled missing values in numerical column 'Rainfall' with the mean:
2.3609181499166656
Filled missing values in numerical column 'Evaporation' with the mean:
5.468231522922462
Filled missing values in numerical column 'Sunshine' with the mean:
7.6111775206611565
Filled missing values in numerical column 'WindGustSpeed' with the mean:
40.03523007167319
Filled missing values in numerical column 'WindSpeed9am' with the mean:
14.043425914971502
Filled missing values in numerical column 'WindSpeed3pm' with the mean:
18.662656778887342
Filled missing values in numerical column 'Humidity9am' with the mean:
68.88083133761887
Filled missing values in numerical column 'Humidity3pm' with the mean:
51.5391158755046
Filled missing values in numerical column 'Pressure9am' with the mean:
1017.6499397983052
Filled missing values in numerical column 'Pressure3pm' with the mean:
1015.2558888309618
Filled missing values in numerical column 'Cloud9am' with the mean:
4.4474612602152455
Filled missing values in numerical column 'Cloud3pm' with the mean:
4.509930082924903
Filled missing values in numerical column 'Temp9am' with the mean:
16.990631415587398
Filled missing values in numerical column 'Temp3pm' with the mean:
21.68339031800974

```

```
[11]: categorical_columns = df.select_dtypes(include=['object']).columns
for col in categorical_columns:
    if col not in binary_columns and df[col].isnull().sum() > 0: # Exclude Yes/
        ↪No columns already handled
        mode_value = df[col].mode()[0]
        df[col].fillna(mode_value, inplace=True)
        print(f"Filled missing values in categorical column '{col}' with the_
        ↪mode: {mode_value}")
```

```

Filled missing values in categorical column 'WindGustDir' with the mode: W
Filled missing values in categorical column 'WindDir9am' with the mode: N
Filled missing values in categorical column 'WindDir3pm' with the mode: SE

```

```
[35]: output_file_path = 'weatherAUS_updated.csv'

# Save the updated DF to a CSV file
df.to_csv(output_file_path, index=False)

print(f"\dataset saved successfully to '{output_file_path}'")
```

\dataset saved successfully to 'weatherAUS_updated.csv'

4 Outliers

```
[13]: file_path = 'weatherAUS_updated.csv'
try:
    df = pd.read_csv(file_path)
    print(f"File '{file_path}' loaded successfully.")
except FileNotFoundError:
    print(f"Error: The file '{file_path}' was not found. Please check the file path and try again.")
    df = None

total_size = df.size
print("Total number of elements:", total_size)
```


File 'weatherAUS_updated.csv' loaded successfully.
Total number of elements: 3345580

```
[33]: # Only numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

plt.figure(figsize=(15, 10))

# Create boxplots for each numerical column
for i, col in enumerate(numerical_columns):
    plt.subplot(4, 4, i + 1)
    sns.boxplot(data=df, x=col, color='skyblue', fliersize=5)
    plt.title(f'Boxplot of {col}')
    plt.tight_layout()

# Show the plots
plt.show()
```



```
[14]: from scipy.stats import zscore

# Select only numerical columns for Z-score calculation
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

# Initialize a dictionary to store outlier counts
outliers_info = {}

[15]: for col in numerical_columns:
    # Calculate Z-scores for the column
    z_scores = zscore(df[col])

    # Identify outliers (where absolute Z-score > threshold)
    outliers = np.abs(z_scores) > 3

    # Get the count of outliers
    outlier_count = outliers.sum() # Sum of True values gives the number of outliers

    # Store the outlier count in the dictionary
    outliers_info[col] = outlier_count
```

```
[16]: for col, outlier_count in outliers_info.items():
    print(f"Number of outliers in '{col}': {outlier_count}")
```

```
Number of outliers in 'MinTemp': 26
Number of outliers in 'MaxTemp': 346
Number of outliers in 'Rainfall': 2482
Number of outliers in 'Evaporation': 1867
Number of outliers in 'Sunshine': 0
Number of outliers in 'WindGustSpeed': 1717
Number of outliers in 'WindSpeed9am': 1362
Number of outliers in 'WindSpeed3pm': 958
Number of outliers in 'Humidity9am': 585
Number of outliers in 'Humidity3pm': 0
Number of outliers in 'Pressure9am': 772
Number of outliers in 'Pressure3pm': 666
Number of outliers in 'Cloud9am': 0
Number of outliers in 'Cloud3pm': 0
Number of outliers in 'Temp9am': 163
Number of outliers in 'Temp3pm': 460
```

```
[17]: # Select only numerical columns for Z-score calculation
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

# Define the Z-score threshold
z_threshold = 3

# Loop through each numerical column
for col in numerical_columns:
    # Calculate Z-scores for the column
    z_scores = zscore(df[col])

    # Calculate the upper and lower bounds based on Z-threshold
    mean_value = df[col].mean()
    std_dev = df[col].std()
    lower_bound = mean_value - z_threshold * std_dev
    upper_bound = mean_value + z_threshold * std_dev

    # Identify an example outlier (if any exist)
    outlier_indices = df[(z_scores > z_threshold) | (z_scores < -z_threshold)].
    ↪index
    if len(outlier_indices) > 0:
        # Take the first outlier for display purposes
        example_index = outlier_indices[0]
        before_value = df.loc[example_index, col]

        # Cap the outlier
        capped_value = (
```

```

        upper_bound if before_value > upper_bound else lower_bound
    )
# Display the before and after values
print(f"Column: '{col}'")
print(f" - Before Capping (Row {example_index}): {before_value}")
print(f" - After Capping: {capped_value}\n")

# Apply capping to the column
df[col] = np.where(df[col] > upper_bound, upper_bound,
                   np.where(df[col] < lower_bound, lower_bound, df[col]))
else:
    print(f"Column: '{col}' has no outliers.\n")

```

Column: 'MinTemp'
- Before Capping (Row 46923): -8.0
- After Capping: -6.903215228998894

Column: 'MaxTemp'
- Before Capping (Row 68): 44.8
- After Capping: 44.48571958519393

Column: 'Rainfall'
- Before Capping (Row 296): 28.8
- After Capping: 27.508380832510234

Column: 'Evaporation'
- Before Capping (Row 6063): 16.4
- After Capping: 14.952854712019302

Column: 'Sunshine' has no outliers.

Column: 'WindGustSpeed'
- Before Capping (Row 8): 80.0
- After Capping: 79.38998835820152

Column: 'WindSpeed9am'
- Before Capping (Row 3308): 48.0
- After Capping: 40.62660283909846

Column: 'WindSpeed3pm'
- Before Capping (Row 52): 48.0
- After Capping: 44.81239984264633

Column: 'Humidity9am'
- Before Capping (Row 6322): 9.0
- After Capping: 12.316535883650047

Column: 'Humidity3pm' has no outliers.

Column: 'Pressure9am'

- Before Capping (Row 12): 994.3
- After Capping: 997.4645392714654

Column: 'Pressure3pm'

- Before Capping (Row 12): 993.0
- After Capping: 995.2639710641031

Column: 'Cloud9am' has no outliers.

Column: 'Cloud3pm' has no outliers.

Column: 'Temp9am'

- Before Capping (Row 5890): 37.6
- After Capping: 36.33829391075331

Column: 'Temp3pm'

- Before Capping (Row 68): 43.4
- After Capping: 42.23356116599605

```
[18]: print("\nUpdated Summary Statistics for Numerical Columns:")
numerical_summary = df.describe(include=[float, int])
display(numerical_summary)
```

Updated Summary Statistics for Numerical Columns:

	MinTemp	MaxTemp	Rainfall	Evaporation	\
--	---------	---------	----------	-------------	---

count	145460.000000	145460.000000	145460.000000	145460.000000	
-------	---------------	---------------	---------------	---------------	--

mean	12.194085	23.224437	1.969541	5.387934	
------	-----------	-----------	----------	----------	--

std	6.365420	7.074739	5.106232	2.594957	
-----	----------	----------	----------	----------	--

min	-6.903215	1.956977	0.000000	0.000000	
-----	-----------	----------	----------	----------	--

25%	7.700000	18.000000	0.000000	4.000000	
-----	----------	-----------	----------	----------	--

50%	12.100000	22.700000	0.000000	5.468232	
-----	-----------	-----------	----------	----------	--

75%	16.800000	28.200000	1.000000	5.468232	
-----	-----------	-----------	----------	----------	--

max	31.291284	44.485720	27.508381	14.952855	
-----	-----------	-----------	-----------	-----------	--

	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	\
--	----------	---------------	--------------	--------------	---

count	145460.000000	145460.000000	145460.000000	145460.000000	
-------	---------------	---------------	---------------	---------------	--

mean	7.611178	39.942171	13.992645	18.624717	
------	----------	-----------	-----------	-----------	--

std	2.729486	12.774919	8.670736	8.577905	
-----	----------	-----------	----------	----------	--

min	0.000000	6.000000	0.000000	0.000000	
-----	----------	----------	----------	----------	--

25%	7.611178	31.000000	7.000000	13.000000	
-----	----------	-----------	----------	-----------	--

50%	7.611178	39.000000	13.000000	18.662657	
-----	----------	-----------	-----------	-----------	--

75%	8.700000	46.000000	19.000000	24.000000	
max	14.500000	79.389988	40.626603	44.812400	
	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm	\
count	145460.000000	145460.000000	145460.000000	145460.000000	
mean	68.894005	51.539116	1017.663747	1015.266767	
std	18.813284	20.471189	6.674746	6.619300	
min	12.316536	0.000000	997.464539	995.263971	
25%	57.000000	37.000000	1013.500000	1011.100000	
50%	69.000000	51.539116	1017.649940	1015.255889	
75%	83.000000	65.000000	1021.800000	1019.400000	
max	100.000000	100.000000	1037.835340	1035.247807	
	Cloud9am	Cloud3pm	Temp9am	Temp3pm	
count	145460.000000	145460.000000	145460.000000	145460.000000	
mean	4.447461	4.509930	16.991338	21.686191	
std	2.265604	2.092954	6.444951	6.833782	
min	0.000000	0.000000	-2.357031	1.133219	
25%	3.000000	4.000000	12.300000	16.700000	
50%	4.447461	4.509930	16.800000	21.400000	
75%	6.000000	6.000000	21.500000	26.200000	
max	9.000000	9.000000	36.338294	42.233561	

5 Visualization and EDA

```
[32]: max_samples = 20000 # Adjust based on memory constraints (Actually so we wont
    ↪take all day)
if len(df) > max_samples:
    df_sampled = df.sample(max_samples, random_state=42)
    print(f"Using a random sample of {max_samples} rows out of {len(df)}.")
else:
    df_sampled = df.copy()
    print(f"Using the full dataset of {len(df)} rows.")

# Separate numerical and categorical columns
numerical_columns = df_sampled.select_dtypes(include=['float64', 'int64']).columns
categorical_columns = df_sampled.select_dtypes(include=['object']).columns

# Standardize numerical data
scaler = StandardScaler()
numerical_data_scaled = scaler.fit_transform(df_sampled[numerical_columns])

# Encode categorical data using sparse matrices
encoder = OneHotEncoder(sparse_output=True, handle_unknown='ignore')
categorical_data_encoded = encoder.
    ↪fit_transform(df_sampled[categorical_columns])
```

```

# Combine numerical and categorical data
if len(categorical_columns) > 0:
    combined_data = hstack([numerical_data_scaled, categorical_data_encoded])
else:
    combined_data = numerical_data_scaled


# Apply PCA
pca = PCA(n_components=2)
pca_embedding = pca.fit_transform(combined_data.toarray()) # Convert sparse matrix to dense if applicable

# K-Means clustering
n_clusters = 5 # Adjust based on your dataset
kmeans = KMeans(n_clusters=n_clusters, random_state=42, n_init=10)
cluster_labels = kmeans.fit_predict(pca_embedding)

plt.figure(figsize=(10, 7))
scatter = plt.scatter(
    pca_embedding[:, 0], pca_embedding[:, 1], c=cluster_labels, cmap='tab10', alpha=0.7, s=10
)
plt.title(f'PCA Visualization with {n_clusters} Clusters')
plt.xlabel('PCA Dimension 1')
plt.ylabel('PCA Dimension 2')
plt.colorbar(scatter, label='Cluster')
plt.grid(True)
plt.show()

```

Using a random sample of 20000 rows out of 145460.



Explained variance by each component: [0.25609104 0.15842998]
Cumulative explained variance: [0.25609104 0.41452102]

```
[20]: max_samples = 20000 # Adjust based on memory constraints
if len(df) > max_samples:
    df_sampled = df.sample(max_samples, random_state=42)
    print(f"Using a random sample of {max_samples} rows out of {len(df)}.")
else:
    df_sampled = df.copy()
    print(f"Using the full dataset of {len(df)} rows.")
```

Using a random sample of 20000 rows out of 145460.

```
[21]: # Preprocess the data for UMAP
# Separate numerical and categorical columns
numerical_columns = df_sampled.select_dtypes(include=['float64', 'int64']).columns
categorical_columns = df_sampled.select_dtypes(include=['object']).columns

# Scale numerical data
```

```

scaler = StandardScaler()
numerical_data_scaled = scaler.fit_transform(df_sampled[numerical_columns])

# Encode categorical data using sparse matrices
encoder = OneHotEncoder(sparse_output=True, handle_unknown='ignore')
categorical_data_encoded = encoder.
    ↪fit_transform(df_sampled[categorical_columns])

```

```

[22]: # Combine numerical and encoded categorical data
if categorical_columns.any():
    combined_data = hstack([numerical_data_scaled, categorical_data_encoded])
else:
    combined_data = numerical_data_scaled

# Suppress the UMAP warning for n_jobs and use random_state for reproducibility
with warnings.catch_warnings():
    warnings.filterwarnings("ignore", category=UserWarning, module="umap")

# Apply UMAP
umap_reducer = umap.UMAP(n_neighbors=15, min_dist=0.1, random_state=42, ↴
    ↪low_memory=True)
embedding = umap_reducer.fit_transform(combined_data)


n_clusters = 5
kmeans = KMeans(n_clusters=n_clusters, random_state=42, n_init=10)
cluster_labels = kmeans.fit_predict(embedding)

```

```

[23]: plt.figure(figsize=(10, 7))
scatter = plt.scatter(embedding[:, 0], embedding[:, 1], c=cluster_labels, ↴
    ↪cmap='tab10', alpha=0.7, s=10)
plt.title(f'UMAP Visualization with {n_clusters} Clusters')
plt.xlabel('UMAP Dimension 1')
plt.ylabel('UMAP Dimension 2')
plt.colorbar(scatter, label='Cluster')
plt.grid(True)
plt.show()


```



```
[24]: # Plot average Rainfall over time
df['Date'] = pd.to_datetime(df['Date'])


rainfall_over_time = df.groupby('Date')['Rainfall'].mean()

plt.figure(figsize=(12, 6))
plt.plot(rainfall_over_time, label='Average Rainfall', color='blue')
plt.title('Rainfall Over Time')
plt.xlabel('Date')
plt.ylabel('Average Rainfall')
plt.legend()
plt.grid(True)
plt.show()
```




```
[25]: # RainTomorrow Distribution
plt.subplot(1, 2, 2)
sns.countplot(data=df, x='RainTomorrow', color='salmon', edgecolor='black') # Using Simplified palette
plt.title('RainTomorrow Distribution')

plt.tight_layout()
plt.show()
```




```
[26]: # RainToday RainTomorrow
      ↪ Distribution plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
sns.countplot(data=df, x='RainToday', color='skyblue', edgecolor='black') # ↪ Simplified palette
plt.title('RainToday Distribution')
```

```
[26]: Text(0.5, 1.0, 'RainToday Distribution')
```




```
[27]: # Correlation heatmap
correlation_matrix = df[numerical_columns].corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Heatmap")
plt.show()
```



```
[28]: # Plot histograms for numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

plt.figure(figsize=(15, 10))
for i, col in enumerate(numerical_columns):
    plt.subplot(4, 4, i+1)
    df[col].hist(bins=30, alpha=0.7, color='skyblue', edgecolor='black')
    plt.title(col)
    plt.tight_layout()
plt.show()
```



6 Machine Learning

```
[29]: from pyspark.sql.functions import col, when

# Initialize Spark session
spark = SparkSession.builder.appName("Random Forest WeatherPrediction").
    getOrCreate()

# Convert df to spark df
spark_df = spark.createDataFrame(df)

# Encode categorical variables into numeric
indexer = StringIndexer(inputCol="Location", outputCol="LocationIndex")
spark_df = indexer.fit(spark_df).transform(spark_df)

indexer_wind_dir = StringIndexer(inputCol="WindGustDir", outputCol="WindGustDirIndex")
spark_df = indexer_wind_dir.fit(spark_df).transform(spark_df)

indexer_rain_today = StringIndexer(inputCol="RainToday", outputCol="RainTodayIndex")
spark_df = indexer_rain_today.fit(spark_df).transform(spark_df)
```

```

# Assemble feature columns into a single vector
assembler = VectorAssembler(
    inputCols=[
        'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
        'WindGustSpeed', 'Humidity9am', 'Humidity3pm', 'Pressure9am', □
        ↪'Pressure3pm',
        'LocationIndex', 'WindGustDirIndex', 'RainTodayIndex'
    ],
    outputCol="features"
)
spark_df = assembler.transform(spark_df)

# Encode target variable as a numeric label
indexer_target = StringIndexer(inputCol="RainTomorrow", outputCol="label")
spark_df = indexer_target.fit(spark_df).transform(spark_df)

# Split data into training and testing sets
train_df, test_df = spark_df.randomSplit([0.8, 0.2], seed=1234)

# Train Random Forest
rf = RandomForestClassifier(labelCol="label", featuresCol="features", □
    ↪numTrees=10, maxDepth=5, maxBins=100)
model = rf.fit(train_df)

# Generate predictions on test data
predictions = model.transform(test_df)

# Map numeric predictions back to 'Yes' or 'No'
predictions = predictions.withColumn(
    "predicted_outcome",
    when(col("prediction") == 0, "No").otherwise("Yes")
)

# Display a random sample of predictions
predictions.sample(withReplacement=False, fraction=0.1, seed=1234).
    ↪select("Date", "Location", "predicted_outcome").show(20)

# Evaluate the model using AUC
evaluator = BinaryClassificationEvaluator(labelCol="label", □
    ↪rawPredictionCol="prediction")
auc = evaluator.evaluate(predictions)
print(f"Area Under ROC (AUC): {auc}")

# Stop Spark session
spark.stop()

```

Date	Location	predicted_outcome
2009-01-28 00:00:00	BadgerysCreek	No
2009-02-02 00:00:00	Cobar	No
2009-02-10 00:00:00	BadgerysCreek	Yes
2009-02-19 00:00:00	BadgerysCreek	No
2009-02-25 00:00:00	Albury	No
2009-02-28 00:00:00	BadgerysCreek	No
2009-03-12 00:00:00	Cobar	No
2009-03-17 00:00:00	Cobar	No
2009-05-04 00:00:00	BadgerysCreek	No
2009-05-04 00:00:00	Cobar	No
2009-05-10 00:00:00	Albury	No
2009-05-13 00:00:00	Albury	No
2009-07-01 00:00:00	Cobar	No
2009-07-19 00:00:00	BadgerysCreek	No
2009-07-23 00:00:00	Cobar	No
2009-08-03 00:00:00	BadgerysCreek	No
2009-08-04 00:00:00	Cobar	No
2009-09-01 00:00:00	Albury	No
2009-09-20 00:00:00	BadgerysCreek	No
2009-10-04 00:00:00	Cobar	No

only showing top 20 rows

Area Under ROC (AUC): 0.6605941347718546

```
[30]: from pyspark.sql.functions import col, when

# Init Spark session
spark = SparkSession.builder.appName("GBT WeatherPrediction").getOrCreate()

# Convert Pandas DataFrame to PySpark DataFrame
spark_df = spark.createDataFrame(df)

# Handle categorical variables
indexer = StringIndexer(inputCol="Location", outputCol="LocationIndex")
spark_df = indexer.fit(spark_df).transform(spark_df)

indexer_wind_dir = StringIndexer(inputCol="WindGustDir", outputCol="WindGustDirIndex")
spark_df = indexer_wind_dir.fit(spark_df).transform(spark_df)

indexer_rain_today = StringIndexer(inputCol="RainToday", outputCol="RainTodayIndex")
spark_df = indexer_rain_today.fit(spark_df).transform(spark_df)
```

```

# feature vector
assembler = VectorAssembler(
    inputCols=[
        'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
        'WindGustSpeed', 'Humidity9am', 'Humidity3pm', 'Pressure9am', □
        ↵'Pressure3pm',
        'LocationIndex', 'WindGustDirIndex', 'RainTodayIndex'
    ],
    outputCol="features"
)
spark_df = assembler.transform(spark_df)

# Encode the target variable 'RainTomorrow'
indexer_target = StringIndexer(inputCol="RainTomorrow", outputCol="label")
spark_df = indexer_target.fit(spark_df).transform(spark_df)

# Split the data into training and testing sets
train_df, test_df = spark_df.randomSplit([0.8, 0.2], seed=1234)

# GBT
gbt = GBTCClassifier(labelCol="label", featuresCol="features", maxIter=10, □
    ↵maxBins=100)

# Fit the model
model = gbt.fit(train_df)

# Make predictions
predictions = model.transform(test_df)

# Map the numeric predictions back to string labels ('Yes' or 'No')
predictions = predictions.withColumn(
    "predicted_outcome",
    when(col("prediction") == 0, "No").otherwise("Yes")
)

# Show predictions
predictions.select("Date", "Location", "predicted_outcome").show()

# Evaluate the model using AUC
evaluator = BinaryClassificationEvaluator(labelCol="label", □
    ↵rawPredictionCol="prediction")
auc = evaluator.evaluate(predictions)
print(f"Area under ROC (AUC): {auc}")

# Stop session
spark.stop()

```

Date	Location	predicted_outcome
2008-12-02 00:00:00	Albury	No
2008-12-05 00:00:00	Albury	No
2008-12-22 00:00:00	Albury	No
2008-12-25 00:00:00	Albury	No
2008-12-26 00:00:00	Albury	No
2008-12-27 00:00:00	Albury	No
2008-12-28 00:00:00	Albury	No
2009-01-02 00:00:00	Cobar	No
2009-01-05 00:00:00	Cobar	No
2009-01-06 00:00:00	Badgerys Creek	No
2009-01-08 00:00:00	Albury	No
2009-01-09 00:00:00	Badgerys Creek	No
2009-01-09 00:00:00	Cobar	No
2009-01-10 00:00:00	Albury	No
2009-01-11 00:00:00	Cobar	No
2009-01-12 00:00:00	Cobar	No
2009-01-14 00:00:00	Badgerys Creek	No
2009-01-14 00:00:00	Cobar	No
2009-01-16 00:00:00	Albury	No
2009-01-18 00:00:00	Cobar	No

only showing top 20 rows

Area under ROC (AUC): 0.7034931086786104

[31]: `from pyspark.sql.functions import col, when`

```
# Initialize Spark session
spark = SparkSession.builder.appName("Logistic Regression WeatherPrediction")\
    .getOrCreate()

# Convert Pandas DataFrame to PySpark DataFrame
spark_df = spark.createDataFrame(df)

# Handle categorical variables by encoding them into numeric indices
indexer_location = StringIndexer(inputCol="Location", outputCol="LocationIndex")
spark_df = indexer_location.fit(spark_df).transform(spark_df)

indexer_wind_dir = StringIndexer(inputCol="WindGustDir", \
    outputCol="WindGustDirIndex")
spark_df = indexer_wind_dir.fit(spark_df).transform(spark_df)
```

```

indexer_rain_today = StringIndexer(inputCol="RainToday", outputCol="RainTodayIndex")
spark_df = indexer_rain_today.fit(spark_df).transform(spark_df)

# Assemble feature columns into a single vector
assembler = VectorAssembler(
    inputCols=[
        'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
        'WindGustSpeed', 'Humidity9am', 'Humidity3pm', 'Pressure9am',
        'Pressure3pm',
        'LocationIndex', 'WindGustDirIndex', 'RainTodayIndex'
    ],
    outputCol="features"
)
spark_df = assembler.transform(spark_df)

# Encode the target variable 'RainTomorrow' as a numeric label
indexer_target = StringIndexer(inputCol="RainTomorrow", outputCol="label")
spark_df = indexer_target.fit(spark_df).transform(spark_df)

# Split the data into training and testing sets
train_df, test_df = spark_df.randomSplit([0.8, 0.2], seed=1234)

# Train the Logistic Regression model
logistic_reg = LogisticRegression(labelCol="label", featuresCol="features",
    maxIter=10, regParam=0.01, elasticNetParam=0.0)
logistic_model = logistic_reg.fit(train_df)

# Make predictions on the test data
predictions = logistic_model.transform(test_df)

# Map the numeric predictions back to string labels ('Yes' or 'No')
predictions = predictions.withColumn(
    "predicted_outcome",
    when(col("prediction") == 0, "No").otherwise("Yes")
)

# Display predictions
predictions.select("Date", "Location", "predicted_outcome").show(20)

# Evaluate the model using AUC
evaluator = BinaryClassificationEvaluator(labelCol="label",
    rawPredictionCol="rawPrediction")
auc = evaluator.evaluate(predictions)
print(f"Area Under ROC (AUC) for Logistic Regression: {auc}")

# Stop Spark session

```

```
spark.stop()
```

Date	Location	predicted_outcome
2008-12-02 00:00:00	Albury	No
2008-12-05 00:00:00	Albury	No
2008-12-22 00:00:00	Albury	No
2008-12-25 00:00:00	Albury	No
2008-12-26 00:00:00	Albury	No
2008-12-27 00:00:00	Albury	No
2008-12-28 00:00:00	Albury	No
2009-01-02 00:00:00	Cobar	No
2009-01-05 00:00:00	Cobar	No
2009-01-06 00:00:00	BadgerysCreek	No
2009-01-08 00:00:00	Albury	No
2009-01-09 00:00:00	BadgerysCreek	No
2009-01-09 00:00:00	Cobar	No
2009-01-10 00:00:00	Albury	No
2009-01-11 00:00:00	Cobar	No
2009-01-12 00:00:00	Cobar	No
2009-01-14 00:00:00	BadgerysCreek	No
2009-01-14 00:00:00	Cobar	No
2009-01-16 00:00:00	Albury	No
2009-01-18 00:00:00	Cobar	No

only showing top 20 rows

Area Under ROC (AUC) for Logistic Regression: 0.8511685455881767

```
[ ]:
```